

Exercise 7

Unstructured Spark exercise

Prior Knowledge

Unix Command Line Shell

Simple Python

Spark Python

Simple SQL syntax

Learning Objectives

Pulling together your skills from previous exercises

Software Requirements

(see separate document for installation of these)

- Apache Spark 2.1.1
- Python 2.7.12
- Jupyter Notebook

Aim

There is a file on your VM that contains some data about health practices (e.g. GP surgeries) in the UK:

`~/datafiles/practices/ukpractices2015.csv`

The CSV file has a header line with titles of each column.

The aim is simple:

I'd like you to calculate the number of practices per postcode prefix for the data. The postcode prefix I define as the first few characters of the postcode up to the space.

Please tell me the number of surgeries for the postcode areas: OX1, SW11.

We are going to do this locally, NOT on EC2.

There are some hints overleaf.

Hints:

1. Create a new Jupyter Notebook as in Exercise 5
2. Use the same databricks CSV reader from Exercise 6 to load the data in
3. You should know enough to do this:
 - a. either as a set of Map/ReduceByKey operations. You could also look at countByKey
 - b. Alternatively, you can do this all in SQL if you like SQL.
4. If you like to mix and match SQL and Map/Reduce you can do that too. I've shown you how to do DataFrame → RDD. The following page shows you how to do RDD → DataFrame:

<https://spark.apache.org/docs/latest/sql-programming-guide.html#interoperating-with-rdds>

5. Ask me or David if you get stuck.