

Exercise 7

More Apache Spark and Python, EC2

Prior Knowledge

Unix Command Line Shell

Simple Python

Learning Objectives

Using Spark on EC2

Accessing S3 files on Spark

Reading CSV files in Spark

Seeing the differences between Spark and Hadoop by performing the Word

Analysis in Spark

Spark SQL

Software Requirements

(see separate document for installation of these)

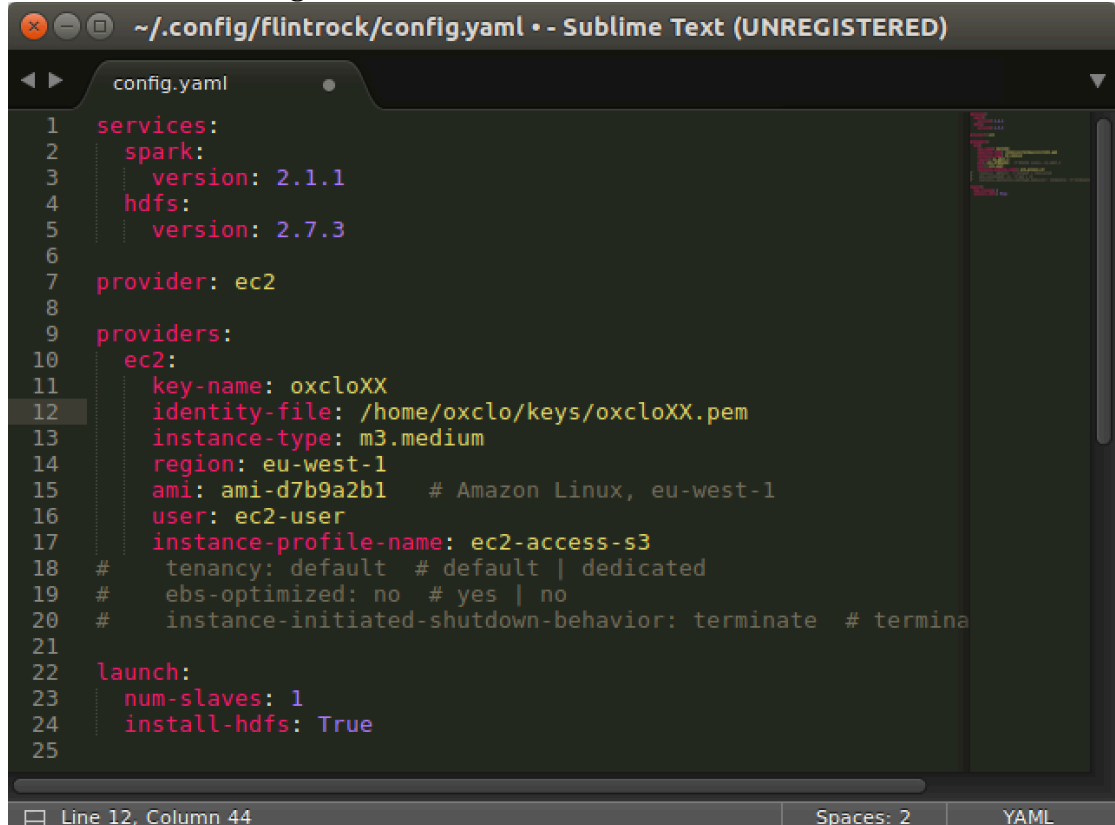
- EC2 credentials
- Flintrock
- Livy

Part A. Starting Spark in EC2

1. Do you remember the Access Key and Secret Key from Exercise 1? You need those now. If you have lost them, ask for help.
2. In a terminal window type:
`export AWS_ACCESS_KEY_ID=<your access key here>`
`export AWS_SECRET_ACCESS_KEY=<your secret key here>`
3. Further commands in this lab need to be in this same command window, because if you start a new window, these environment variables won't be there. In a more serious setup you would either add these to your profile or put them in a command-shell program so you could call them easily.
4. There is a project from the creators of Spark to run it in EC2, but it is not very good! Instead we will use a tool called **flintrock**
5. Before we can use flintrock, you need to modify the config file for flintrock so that it uses your own keys. Edit the flintrock config file:

```
subl ~/.config/flintrock/config.yaml
```

It will look something like:



```
1  services:
2    spark:
3      version: 2.1.1
4    hdfs:
5      version: 2.7.3
6
7  provider: ec2
8
9  providers:
10   ec2:
11     key-name: oxcloXX
12     identity-file: /home/oxclo/keys/oxcloXX.pem
13     instance-type: m3.medium
14     region: eu-west-1
15     ami: ami-d7b9a2b1 # Amazon Linux, eu-west-1
16     user: ec2-user
17     instance-profile-name: ec2-access-s3
18     # tenancy: default # default | dedicated
19     # ebs-optimized: no # yes | no
20     # instance-initiated-shutdown-behavior: terminate # termina
21
22   launch:
23     num-slaves: 1
24     install-hdfs: True
25
```

The source for this is here:

<https://freo.me/flintrock-conf>

This is modified in a couple of ways. Firstly, it gives the Ireland region and AMI files. Secondly, there is an “instance-profile-name”. This is a AWS feature that gives the running VM access to other APIs - in this case S3.

6. Change the key name and identity file to match your key name and identity file. Save the file.
7. You should now be able to launch a cluster in Amazon:

```
cd flintrock
./flintrock launch oxcloXX-sc

(using your XX)
```

8. Now you should see something like:

```

Launching 2 instances...
[54.154.17.100] SSH online.
[54.154.17.100] Configuring ephemeral storage...
[54.154.17.100] Installing Java 1.8...
[34.253.201.139] SSH online.
[34.253.201.139] Configuring ephemeral storage...
[34.253.201.139] Installing Java 1.8...
[54.154.17.100] Installing HDFS...
[34.253.201.139] Installing HDFS...
[54.154.17.100] Installing Spark...
[34.253.201.139] Installing Spark...
[34.253.201.139] Configuring HDFS master...
[34.253.201.139] Configuring Spark master...
HDFS online.
Spark Health Report:
  * Master: ALIVE
  * Workers: 1
  * Cores: 1
  * Memory: 2.7 GB
launch finished in 0:03:49.

```

9. Let's login to the master (all one line):

```
./flintrock login oxcloXX-sc
```

You see something like:

```
Warning: Permanently added '34.253.201.139' (ECDSA) to the list
of known hosts.
Last login: Mon Jul 10 18:55:35 2017 from host109-156-251-
208.range109-156.btcentralplus.com
```

```

_ | ( _ /
_ | \ _ | _ |

```

Amazon Linux AMI

```
https://aws.amazon.com/amazon-linux-ami/2017.03-release-notes/
1 package(s) needed for security, out of 1 available
Run "sudo yum update" to apply all updates.
[ec2-user@ip-172-31-6-32 ~]$
```

10. This basically just SSH's you into the master. You could do the same from the EC2 console as before.

11. Now start pyspark once again.

This time we are going to add in a Spark Package that supports accessing S3 data (Amazon object storage). Once again, all one line

```
pyspark --master spark://0.0.0.0:7077
--packages com.amazonaws:aws-java-sdk-
pom:1.10.34,org.apache.hadoop:hadoop-aws:2.7.2
```


16. Now let's load the data into a DataFrame. (one line)

```
df = sqlContext.read.format('com.databricks.spark.csv').
options(header='true', inferschema='true').
load('s3a://oxclo-wind/2015/*')
```

17. Spark should go away and think a bit, and also show some ephemeral log lines about the staging.

18. The df object we have is not an RDD, but instead a DataFrame. This is basically a SQL construct. But we can easily convert it into an RDD.

19. We can print a nice table showing the first few rows with:

```
df.show(4)
```

Station_ID	Station_Name	Location_Label	Interval_Minutes	Interval_End_Time	Wind_Velocity_Mtr_Sec	Wind_Direction_Variance_Deg	Wind_Direction_Deg	Ambient_Temperature_Deg_C	Global_Horizontal_Irradiance
SF15	Warnerville Switc...	Warnerville	5	2015-01-57 00:05	1.628	8.1	148.5	0.92	0.861
SF15	Warnerville Switc...	Warnerville	5	2015-01-57 00:10	1.519	9.4	151.1	0.717	0.864
SF15	Warnerville Switc...	Warnerville	5	2015-01-57 00:15	1.482	8.7	142.7	0.627	0.859
SF15	Warnerville Switc...	Warnerville	5	2015-01-57 00:20	1.085	6.895	141.8	0.5	0.862

only showing top 4 rows

(I shrunk this so you can see the table nicely!)

20. We can also convert the DataFrame into an RDD, allowing us to do functional programming on it (map/reduce/etc)

```
winds = df.rdd
```

21. Let's do the normal step of mapping the data into a simple <K,V> pair.
Each column in the row can be accessed by the syntax e.g. row.Station_ID

We can therefore map our RDD with the following:

```
mapped = winds.map(lambda s: (s.Station_ID, s.Wind_Velocity_Mtr_Sec))
```

22. We can simply calculate the maximum values with this reducer:

```
maxes = mapped.reduceByKey(lambda a, b: a if (a>b) else b)
```

23. And once again collect / print:

```
for (k,v) in maxes.collect(): print k,v
```

24. You will see a bunch of log before the following appears:

```
SF18 10.57
SF36 11.05
SF37 7.079
SF15 7.92
SF04 34.12
SF17 5.767
```

25. You can also turn the response of a collect into a Python Map, which is handy. Try this:

```
maxes.collectAsMap()['SF04']
```

26. You can also try:
`print maxes.collectAsMap()`

PART B – Using SQL

27. There is an easier way to do all this if you are willing to write some SQL.

28. First we need to give our DataFrame a table name:

```
df.registerTempTable('wind')
```

29. Now we can use a simple SQL statement against our data.

ALL ON ONE Line type:

```
sqlContext.sql("SELECT Station_ID, avg(Wind_Velocity_Mtr_Sec) as  
avg,max(Wind_Velocity_Mtr_Sec) as max from wind group by  
Station_ID").show()
```

30. Bingo you should see a lot of log followed by:

Station_ID	avg	max
SF36	2.464172530911313	11.05
SF37	2.260403505500663	7.079
SF04	2.300981748124102	34.12
SF15	1.8214145677504483	7.92
SF17	0.5183500253485376	5.767
SF18	2.2202234391695437	10.57


31. Recap. So far we have:

- Started Spark in EC2
- Loaded data from S3
- Used SQL to read in CSV files
- Explored Map/Reduce on those CSV files
- Used SQL to query the data.

32. Find the IP address of the Spark Master. There are two ways. Firstly, it showed up in the console when you first launched the flintrock cluster: [34.253.201.139] Configuring Spark master...

Alternatively, you can find it as "oxcloXX-sc-master" in the EC2 instances.

33. Go to <http://xx.xx.xx.xx:8080> using the master's IP address. You should see something like:

 **Spark Master at spark://ec2-34-253-201-139.eu-west-1.compute.amazonaws.com:7077**

URL: spark://ec2-34-253-201-139.eu-west-1.compute.amazonaws.com:7077
 REST URL: spark://ec2-34-253-201-139.eu-west-1.compute.amazonaws.com:6066 (cluster mode)
 Alive Workers: 1
 Cores in use: 1 Total, 1 Used
 Memory in use: 2.7 GB Total, 1024.0 MB Used
 Applications: 1 Running, 0 Completed
 Drivers: 0 Running, 0 Completed
 Status: ALIVE

Workers

Worker Id	Address	State	Cores	Memory
worker-20170710185543-172.31.1.109-39733	172.31.1.109:39733	ALIVE	1 (1 Used)	2.7 GB (1024.0 MB Used)

Running Applications

Application ID	Name	Cores	Memory per Node	Submitted Time	User	State	Duration
app-20170710192533-0000	(kill) PySparkShell	1	1024.0 MB	2017/07/10 19:25:33	ec2-user	RUNNING	10 s

Completed Applications

Application ID	Name	Cores	Memory per Node	Submitted Time	User	State	Duration
----------------	------	-------	-----------------	----------------	------	-------	----------

34. Quit the pyspark shell:
quit()

35. Exit the SSH session:
exit

36. If you want you can try adding another slave and then rerun the analysis. You can see the extra core working in the Web UI

`./flintrock add-slave --num-slaves 1 oxcloXX-sc`

To save you retyping all that spark code, look here:

<https://freo.me/wind-sql>

37. If you are planning to do the **Jupyter on EC2 exercise** straight away, then you can start it now and use your existing flintrock/EC2 cluster. Otherwise please follow the next instruction to shut down the EC2 instances.

38. We must remember to stop our cluster as well (its costing money...) From Ubuntu terminal where you started the Spark cluster

`./flintrock destroy oxcloXX-sc`

Type y when prompted.

39. Congratulations, this lab is complete.